

# SQL: A COMMERCIAL DATABASE LANGUAGE

Complex Queries,  
Aggregate Functions and Grouping

# Outline of Chapters 8, 9

---

1. Introduction
2. Data Definition, Basic Constraints, and Schema Changes
3. Basic Queries
4. More complex Queries
5. Aggregate Functions and Grouping
6. Summary of SQL queries
7. Data Change statements
8. Views
9. Complex Constraints
10. Database Programming

## 4. More complex queries

---

- Because of the generality and expressive power of SQL , there are many additional features (e.g. query nesting, and grouping/aggregation) that allow users to specify more complex queries.

## 4.1 Nested queries (1)

---

- A complete SELECT query, called a *nested query*, can be specified within the WHERE-clause of another query, called the *outer query*.

- Example:

**Query 1:** Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT  FNAME, LNAME, ADDRESS
FROM    EMPLOYEE
WHERE   DNO   IN   (SELECT      DNUMBER
                     FROM        DEPARTMENT
                     WHERE        DNAME='Research');
```

- The *nested query* selects the number of the 'Research' department.
- The *outer query* selects an EMPLOYEE tuple if its DNO value is in the result of the nested query.
- The *comparison operator IN* compares a value v with a set (or multi-set) of values V, and evaluates to TRUE if v is one of the elements in V.

## 4.1 Nested queries (2)

---

- Example:

**Query 4:** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT  DISTINCT PNUMBER
FROM      PROJECT
WHERE     PNUMBER IN
          (SELECT  PNUMBER
FROM          PROJECT, DEPARTMENT, EMPLOYEE
WHERE        DNUM=DNUMBER AND MGRSSN=SSN
              AND  LNAME='Smith')

OR

          PNUMBER IN
          (SELECT  PNO
FROM          WORKS_ON, EMPLOYEE
WHERE        ESSN=SSN AND LNAME='Smith');
```

## 4.1 Nested queries (3)

---

- The IN comparison operator can also compare a tuple of values in parentheses with a set of *union-compatible* tuples.
- Example:

```
SELECT  DISTINCT ESSN
FROM    WORKS_ON
WHERE   (PNO, HOURS) IN
        (SELECT PNO, HOURS
         FROM    WORKS_ON
         WHERE   ESSN = '123456789');
```

Select the social security number of all employees who work the same (project, hours) combination on a project that the employee whose SSN is '123456789' works on.

## 4.1 Nested queries (4)

---

- Other comparison operators can also be used to *compare a single value  $v$  (typically an attribute name) to a set or multiset  $V$  (typically a nested query)*.
- $\theta$  **ANY** (  $\theta$  **SOME** ) and  $\theta$  **ALL**  
where  $\theta$  is one of  $=, <, <=, >, >=, <>$ .  
 $=$  **ANY** (or  $=$  **SOME**) is equivalent to **IN**.
- Example: Retrieve the names of the employees whose salary is greater than the salary of all the employees in department number 5.

```
SELECT  LNAME, FNAME
FROM    EMPLOYEE
WHERE   SALARY > ALL (SELECT SALARY
                        FROM    EMPLOYEE
                        WHERE    DNO = 5);
```

## 4.1 Nested queries (5)

---

- In general, we can have *several levels of nested queries*.
- Ambiguity can arise among attribute names if the same attribute exists in relations in the FROM clauses of both the outer query and a nested query.

### Rule:

- A reference to an *unqualified attribute* refers to the relation declared in the *innermost nested query*.
- To refer to an attribute of a relation specified in an outer query, we can *specify and refer to an alias for that relation*.

-- examples follow.



## 4.1 Nested queries (6)

---

- Example:

```
SELECT  LNAME, FNAME
FROM    EMPLOYEE
WHERE   SALARY > ALL
        (SELECT  SALARY
         FROM     EMPLOYEE
         WHERE    DNO = 5);
```

Attributes SALARY and DNO in the nested query refer to the EMPLOYEE relation declared in the nested query.

## 4.2 Correlated nested queries (1)

---

- If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query*, the two queries are said to be *correlated*.
- The result of a correlated nested query is *different for each tuple (or combination of tuples) of the relation(s) of the outer query*.
- We can understand a correlated query, by considering that *the nested query is evaluated once for each tuple (or combination of tuples) of the outer query*.

## 4.2 Correlated nested queries (2)

---

- Example:

**Query 16.** Retrieve the name of each employee who has a dependent with the same first name and sex as the employee.

```
SELECT  E.FNAME, E.LNAME
FROM    EMPLOYEE AS E
WHERE    E.SSN IN
            (SELECT ESSN
             FROM    DEPENDENT
             WHERE    SEX = E.SEX AND
                     E.FNAME = DEPENDENT_NAME);
```

We can think that Query 16 is evaluated as follows: for each EMPLOYEE tuple (outer query) evaluate the nested query, which retrieves the ESSN values for all DEPENDENT tuples with the same sex and first name as the EMPLOYEE tuple.

## 4.2 Correlated nested queries (3)

---

- In general, a query with nested SELECT-FROM-WHERE blocks that uses the IN comparison operator can be expressed using a single query block.
- Example:  
Query 16 can be rewritten as follows:

```
SELECT  E.FNAME, E.LNAME
FROM    EMPLOYEE AS E, DEPENDENT AS D
WHERE    D.ESSN = E.SSN AND D.SEX = E.SEX AND
          E.FNAME = D.DEPENDENT_NAME);
```

## 4.3 The EXISTS function (1)

---

- The EXISTS function in SQL is used to check whether the result of a correlated nested query is empty (contains no tuples) or not.
- **Example:** Query 16 can be rewritten as follows:

```
SELECT  E.FNAME, E.LNAME
FROM    EMPLOYEE AS E
WHERE   EXISTS
        (SELECT  *
         FROM    DEPENDENT
         WHERE   E.SSN = ESSN AND E.SEX = SEX AND
                 E.FNAME = DEPENDENT_NAME);
```

- We can think that this query is evaluated as follows: *for each EMPLOYEE tuple (outer query)* the nested query which retrieves all DEPENDENT tuples with the same social security number, sex, and first name as the EMPLOYEE tuple is evaluated. If the *result is not empty*, select the EMPLOYEE tuple.

## 4.3 The EXISTS function (2)

---

- The function NOT EXISTS can also be used.

**EXISTS(Q)** returns **TRUE** if *there is at least one tuple in the result of Q*, and **FALSE** otherwise.

**NOT EXISTS(Q)** returns **TRUE** if *there are no tuples in the result of Q*, and **FALSE** otherwise.

- Example: Query 6. Retrieve the names of employees who have no dependents.

```
SELECT  FNAME, LNAME
FROM    EMPLOYEE
WHERE   NOT EXISTS (SELECT  *
                     FROM    DEPENDENT
                     WHERE   SSN=ESSN) ;
```

## 4.3 The EXISTS function (3)

---

- Another example: Query 7. List the names of managers who have at least one dependent.

```
SELECT  FNAME, LNAME
FROM    EMPLOYEE
WHERE   EXISTS (SELECT  *
                  FROM    DEPENDENT
                  WHERE   SSN = ESSN)
AND
        EXISTS (SELECT  *
                  FROM    DEPARTMENT
                  WHERE   SSN = MGRSSN);
```

Here we specify two nested correlated queries. Can we specify the query using one nested query? No nested queries?

## 4.3 The EXISTS function (\*1)

---

- Another example:

Query 7. List the names of managers who have at least one dependent.

```
SELECT    FNAME, LNAME
FROM      EMPLOYEE
WHERE      EXISTS
              (SELECT *
               FROM    DEPARTMENT, DEPENDENT
               WHERE  SSN = ESSN AND SSN = MGRSSN );
```



## 4.3 The EXISTS function (\*2)

---

- Another example:

Query 7. List the names of managers who have at least one dependent.

```
SELECT    FNAME, LNAME
FROM      EMPLOYEE, DEPARTMENT, DEPENDENT
WHERE      SSN = ESSN AND SSN = MGRSSN;
```

Note that, in contrast to the previous rewritings of Query 7, this one may contain duplicates in the answer even if there are no two employees with the same FNAME and LNAME (Why?)

## 4.3 The EXISTS function (4)

---

- Example (cont.):

Query 3. Retrieve the name of each employee who works on all the projects.

We can express this query by using one level of nesting and the set theoretic operation of set difference.

```
SELECT  FNAME, LNAME  
FROM    EMPLOYEE  
WHERE   NOT EXISTS  
      (SELECT  PNUMBER  
       FROM    PROJECT)  
       EXCEPT  
       (SELECT  PNO  
        FROM    WORKS_ON  
        WHERE   ESSN = SSN) );
```

## 4.3 The EXISTS function (5)

---

- Example (cont.):

Query 3. Retrieve the name of each employee who works on all the projects.

Another option for Query 3 is to use a two-level nesting.

```
SELECT FNAME, LNAME
FROM   EMPLOYEE
WHERE NOT EXISTS
      (SELECT   *
       FROM     PROJECT
       WHERE NOT EXISTS
          (SELECT *
           FROM    WORKS_ON
           WHERE   ESSN = SSN AND PNO = PNUMBER));
```

## 5. Aggregate functions and grouping

---

- Because grouping and aggregation are required in many database applications, SQL has features that incorporate these concepts.
- SQL has a number of built-in aggregate functions: **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**.

The **COUNT** function returns the number of tuples or values as specified in the query.

**SUM**, **MAX**, **MIN**, and **AVG** are applied to a multiset of numeric values and return, respectively, the sum, maximum value, minimum value, and average of those values.

**MIN** and **MAX** can be used with non numeric attributes if their domains have a total ordering (e.g. domain DATE).

## 5.1 Aggregate queries (1)

---

- Example:

**Query 19.** Find the sum of the salaries, the maximum salary, the minimum salary, and the average salary among all employees.

```
SELECT  SUM(SALARY), MAX(SALARY), MIN(SALARY),  
        AVG(SALARY)  
FROM    EMPLOYEE ;
```

Some SQL implementations *may not allow more than one function* in the SELECT-clause

## 5.1 Aggregate queries (2)

---

- Example:

**Query 20.** Find the sum of the salaries, the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department .

```
SELECT  SUM(SALARY), MAX(SALARY), MIN(SALARY),  
          AVG(SALARY)  
FROM    EMPLOYEE, DEPARTMENT  
WHERE   DNO = DNUMBER AND DNAME='Research' ;
```

## 5.1 Aggregate queries (3)

---

- Example:

**Queries 21 and 22.** Retrieve the total number of employees in the company (Q21), and the number of employees in the 'Research' department (Q22) .

```
SELECT  COUNT (*)  
FROM    EMPLOYEE;
```

```
SELECT  COUNT (*)  
FROM    EMPLOYEE, DEPARTMENT  
WHERE   DNO = DNUMBER AND DNAME = 'Research';
```

**COUNT(\*)** returns the number of rows (tuples) in the result of the query.

## 5.1 Aggregate queries (4)

---

- The COUNT function can also be used to count values in a column rather than tuples.
- Example:  
**Query 23.** Count the number of distinct salary values in the database.

```
SELECT  COUNT ( DISTINCT SALARY)
FROM    EMPLOYEE;
```

If **COUNT**(SALARY) had been used instead of **COUNT ( DISTINCT SALARY)**, all the **non-null** SALARY values in the EMPLOYEE relation (**including duplicates**) would have been counted.



## 5.1 Aggregate queries (5)

---

- Aggregate queries are one-tuple queries. They can be nested in the WHERE clause of an outer query to retrieve a summary value from the database.
- Example:  
**Query 5.** Retrieve the names of all employees who have two or more dependents.

```
SELECT  LNAME, FNAME
FROM    EMPLOYEE
WHERE    (SELECT COUNT(*)
           FROM DEPENDENT
           WHERE SSN = ESSN) >= 2;
```

## 5.2 Grouping queries (1)

---

- In many cases, we want to apply the aggregate functions *to subgroups of tuples in a relation*.
- Each subgroup of tuples consists of the set of tuples that have *the same value* for some attribute(s), called *grouping attribute(s)*.
- The aggregate function is applied *to each subgroup independently*.
- SQL has a **GROUP BY**-clause for specifying the grouping attributes.
- The grouping attributes should *also appear in the SELECT-clause*, so that the value resulting from applying each aggregate function to a group of tuples appears along with values of the grouping attribute(s).

## 5.2 Grouping queries (2)

---

- Example:

**Query 24.** For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
SELECT  DNO, COUNT (*), AVG (SALARY)
FROM    EMPLOYEE
GROUP BY DNO;
```

- In Q24, the EMPLOYEE tuples are divided into groups--each group having the same value for the grouping attribute DNO.
- The COUNT and AVG functions are applied to each such group of tuples separately.
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples.

## 5.2 Grouping queries (3)

- Example (cont):

(a.)

FNAME	MINIT	LNAME	SSN	• • •	SALARY	SUPERSSN	DNO
John	B	Smith	123456789	- - -	30000	333445555	5
Franklin		Wong	333445555		40000	888665555	5
Ramesh	K	Narayan	888664444		38000	333445555	5
Jayce	A	English	453453453		25000	333445555	5
Aida	J	Zelaya	999667777		25000	997654321	4
Jennifer	S	Wallace	987654321		43000	888665555	4
Ahmad	V	Jabbar	987987987		25000	997654321	4
James	E	Bong	888665555		55000	null	1

  

DNO	COUNT (*)	AVG (SALARY)
5	4	33250
4	3	31000
1	1	55000

Result of Q24.

Grouping EMPLOYEE tuples by the value of DNO.

...

## 5.2 Grouping queries (4)

---

- A join condition can be used in conjunction with grouping.
- Example:  
**Query 25.** For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT    PNUMBER, PNAME, COUNT (*)  
FROM      PROJECT, WORKS_ON  
WHERE     PNUMBER = PNO  
GROUP BY PNUMBER, PNAME;
```

In this case, the grouping and aggregate functions are applied *after* the joining of the two relations.

## 5.3 The HAVING clause (1)

---

- Sometimes we want to retrieve the values of aggregate functions for only those *groups that satisfy certain conditions*.
- The HAVING-clause is used for specifying a selection condition *on groups* (rather than *on individual tuples*)
- Example:  
*Query 26.* For each project *on which more than two employees work*, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT    PNUMBER, PNAME, COUNT (*)
FROM      PROJECT, WORKS_ON
WHERE     PNUMBER = PNO
GROUP BY  PNUMBER, PNAME
HAVING    COUNT (*) > 2;
```

## 5.3 The HAVING clause (2)

- Example (cont.):

PNAME	PNUMBER		ESSN	PNO	HOURS
ProductX	1	-	123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666664444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10		333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20	-	333445555	20	10.0
Reorganization	20		987854321	20	15.0
Reorganization	20		888665555	20	null
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987854321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by the HAVING condition of Q26.

After applying the where clause but before applying HAVING.

PNAME	PNUMBER		ESSN	PNO	HOURS
ProductY	2	-	123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
Computerization	10		333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987854321	20	15.0
Reorganization	20		888665555	20	null
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987854321	30	20.0
Newbenefits	30		999887777	30	30.0

PNAME	COUNT (*)
ProductY	
Computerization	
Reorganization	
Newbenefits	

Result of Q26  
(PNUMBER not shown).

After applying the HAVING clause condition.

## 5.3 The HAVING clause (3)

---

- Notice that, while *selection conditions in the WHERE-clause limit the tuples to which aggregate functions are applied*, the **HAVING-clause** serves to choose *whole groups*.
- Example:  
**Query 27.** For each project, retrieve the project number, project name, and the number of employees *from department 5* who work on that project.

```
SELECT    PNUMBER, PNAME, COUNT (*)  
FROM      PROJECT, WORKS_ON, EMPLOYEE  
WHERE     PNUMBER = PNO AND SSN =ESSN AND DNO = 5  
GROUP BY PNUMBER, PNAME;
```

Here we restrict tuples in the relation (namely those that have DNO=5) *before the grouping*.



## 5.3 The HAVING clause (4)

---

- We have to be careful when two different conditions apply (one to the tuples in the relation and another to the aggregate function).
- Example:  
**Query 28.** Count the *total* number of employees whose salary exceed \$40000 in each department, but only for departments where more than five employees work.

The following query is INCORRECT (why?):

```
SELECT      DNAME, COUNT (*)  
FROM        DEPARTMENT, EMPLOYEE  
WHERE       DNUMBER = DNO AND SALARY > 40000  
GROUP BY   DNAME  
HAVING      COUNT(*) > 5;
```

## 5.3 The HAVING clause (5)

---

- Example (cont.):

A correct formulation of the query:

```
SELECT    DNAME, COUNT (*)
FROM      DEPARTMENT, EMPLOYEE
WHERE      DNUMBER = DNO AND SALARY > 40000 AND
             DNO IN (SELECT    DNO
                       FROM      EMPLOYEE
                       GROUP BY DNO
                       HAVING    COUNT(*) > 5)
GROUP BY DNAME;
```

We assume that department names are unique (DNAME is a secondary key of DEPARTMENT).

## 6. Summary of SQL queries (1)

---

- A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory. The clauses are specified in the following order

**SELECT** <attribute list>

**FROM** <table list>

[**WHERE** <condition>]

[**GROUP BY** <grouping attribute(s)>]

[**HAVING** <group condition>]

[**ORDER BY** <attribute list>]

- The **SELECT-clause** lists the attributes or functions to be retrieved
- The **FROM-clause** specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The **WHERE-clause** specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- **GROUP BY** specifies grouping attributes.
- **HAVING** specifies a condition for selection of groups
- **ORDER BY** specifies an order for displaying the result of a query.

## 6. Summary of SQL queries (2)

---

- A query is evaluated *conceptually* by applying (in this order)
  - the FROM clause (to compute the cross-product of the tables and joined tables involved in the query).
  - the WHERE clause (to delete rows in the cross-product that fail the conditions in the WHERE clause).
  - the GROUP BY clause (to group rows).
  - the HAVING clause (to eliminate groups that fail the conditions in the HAVING clause).
  - The SELECT clause (to eliminate columns that do not appear in the SELECT list, to eliminate duplicate rows if DISTINCT is specified, and to apply aggregate functions).
  - the ORDER BY clause (to sort the query result).

## 6. Summary of SQL queries (3)

---

- In general, the *conceptual evaluation* is not an efficient way of evaluating a query in a real system.
- A DBMS has special query optimization routines to decide on an *efficient query evaluation plan*.
- In general, there are *different ways* to specify the *same query* in SQL.
- It is the responsibility of the DBMS to execute a query efficiently independently of the way it is specified by the user.
- In practice, however, different specifications of the same query may result in internal evaluations plans with different cost.
- The user should be aware of which types of constructs in a query are more expensive to process than others.